

Lecture 9 – Monte Carlo Simulation

Dr. Jing Li

Department of Computing

The Hong Kong Polytechnic University

22&24 Mar 2022

Looking Back (Lecture 9)

- Simulations with R
 - How to generate random number (from a distribution).
 - Why and how to set random seeds.
 - Random sampling (from a given set of objects)
 - **Example:** simulating a linear model.
- **Case Study:** Analysis of PM 2.5 changes in U.S.
 - How to load data and pick up the data we want
 - How to analyze the data with graphs.
 - Examine the data of the entire U.S.
 - Probe into details of data collected by a monitor.
 - Compare results of different states.

Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- Monte Carlo Simulation with R
 - Estimating π
 - Simulating Product Demand
 - Simulating Coin Flipping
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM

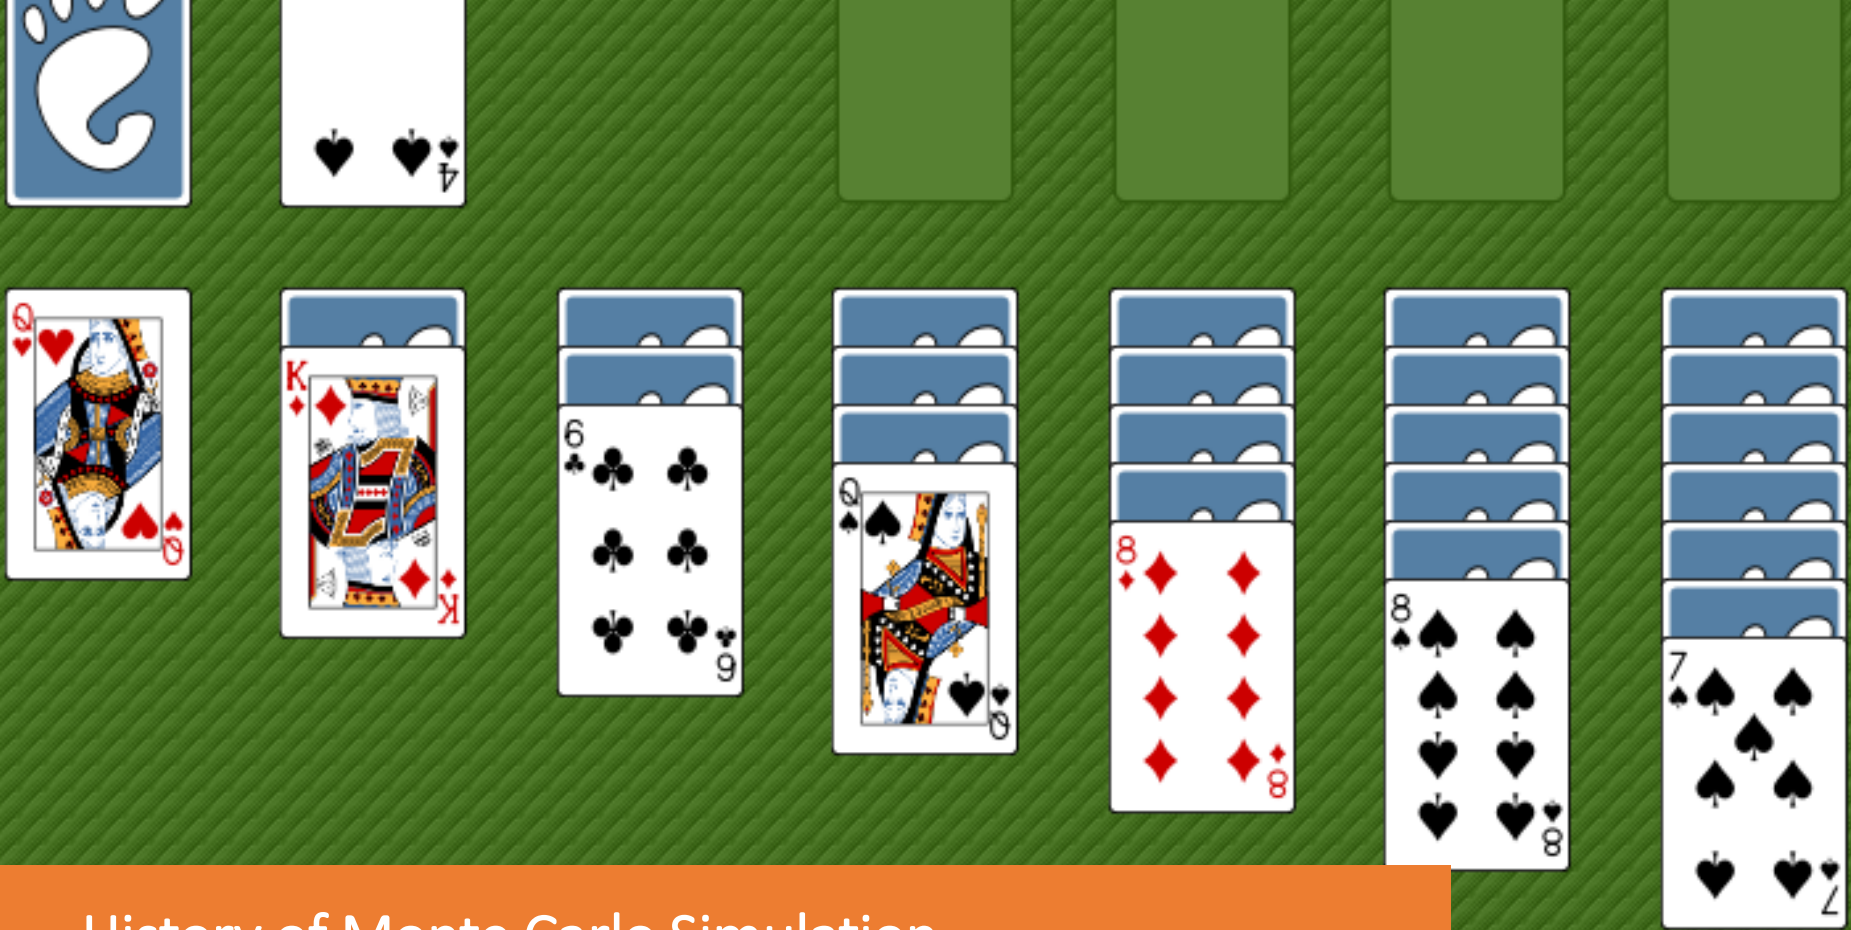
Roadmap

- Introduction: *History and Motivation*
- Key Factors in Monte Carlo Simulation
- Monte Carlo Simulation with R
 - Estimating π
 - Simulating Product Demand
 - Simulating Coin Flipping
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM



5

Where are is it?



History of Monte Carlo Simulation

- *Stanislaw Ulam*, recovering from an illness, was playing a lot of solitaire
- Tried to figure out probability of winning and failed.
- Thought about playing lots of hands and counting number of wins but decided it would take years.
- Asked Von Neumann if he could build a program to simulate many hands on *ENIAC*.

Monte Carlo Simulation

- A method of estimating the value of an unknown quantity using the principles of *inferential statistics*
- Inferential statistics
 - **Population:** a set of examples
 - **Sample:** a proper subset of a population
 - **Key fact:** a *random* sample tends to exhibit the same properties as the population from which it is drawn

Roadmap

- Introduction: *History and Motivation*
- **Key Factors in Monte Carlo Simulation**
- Monte Carlo Simulation with R
 - Estimating π
 - Simulating Product Demand
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM

An example of Coin Flipping

- Given a single coin, estimate fraction of heads you would get if you flipped the coin an infinite number of times
- Consider one flip and heads on!



How confident would you be about answering 1.0?


An example of Coin Flipping

- Given a single coin, estimate fraction of heads you would get if you flipped the coin an infinite number of times
- Consider two flip and both with heads on!



Do you think the next will come up heads?


An example of Coin Flipping



Flip a coin 100 times, all with heads on.

Now do you think that the next flip will come up heads?


An example of Coin Flipping



Flip a coin 100 times, 52 heads and 48 tails.

Do you think that the probability of the next flip coming up heads is $52/100$?

An example of Coin Flipping



Flip a coin 100 times, 52 heads and 48 tails.

Do you think that the probability of the next flip coming up heads is $52/100$?

Your best guess but with low confidence!

Why the Difference in Confidence?

- Confidence in our estimate depends upon two things
 - **Size of sample** (e.g., 100 versus 2)
 - **Variance of sample** (e.g., all heads versus 52 heads)
- As the variance grows, we need larger samples to have the same degree of confidence!



Recall: probability

- *Probability* is a numerical description of *how likely an event is to occur* and or *how likely that a proposition is true*. --- From [Wikipedia](#).
- We run a random experiment n times, during which an event A occurs m times, then we say the *frequency* of A 's occurrence is $f_A = \frac{m}{n}$.
- When n is large enough, f_A will be very close to a value p , which is defined as the probability of A to occur, i.e., $\lim_{n \rightarrow +\infty} f_A \equiv P(A) = p$
 - When we toss a coin, the probability of “heads up” is 0.5

Monte Carlo Principal

- In repeated independent tests with the same actual probability p of a particular outcome in each test, the chance that the fraction of times that outcome occurs differs from p converges to zero as the number of trials goes to infinity
- *Intuition: if deviations from expected behaviour occur, these deviations are likely to be evened out by opposite deviations in the future.*

Law of Large Numbers (different from the one we learned in Lecture 3)

Sampling Space of Possible Outcomes

- Never possible to guarantee perfect accuracy through sampling
- **Key question:**
 - How many samples do we need to look at before we can have justified confidence on our answer?
- Depends upon **variability** in underlying distribution

Roadmap


- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- **Monte Carlo Simulation with R**
 - Estimating π
 - Simulating Product Demand
 - Simulating Coin Flipping
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM

Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- **Monte Carlo Simulation with R**
 - **Estimating π**
 - Simulating Product Demand
 - Simulating Coin Flipping
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM


Estimating π

- What is the ratio of the area of the circle to the area of the square is?
 - $p = \frac{\pi r^2}{(2r)^2} = \frac{\pi}{4}$
- Randomly generate points (x, y)
 - $-1 \leq x, y \leq 1$ (let $r = 1$)
 - Count the number of objects inside the circle and the total number of samples.
 - The ratio of the two counts is an estimate of the ratio of the two areas, which is $\pi/4$.
 - Multiply the result by 4 to estimate π .

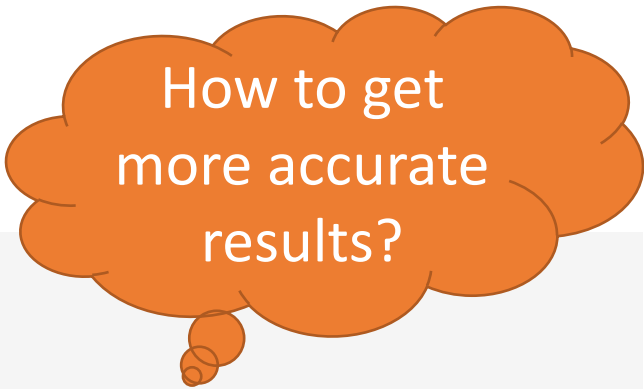


Estimating π

- Generate 1,000 points (x, y) , $-1 \leq x, y \leq 1$.
- Calculate the distance of each point from $(0,0)$ and determine whether each point is inside the circle.
- We may compute the ratio of samples inside the circle and all samples
- Multiply the ratio by 4 to estimate π .



Estimating π



How to get
more accurate
results?

```
no_of_points <- 1000
#runif samples from a uniform distribution
x<- runif(no_of_points, -1, 1)
y<- runif(no_of_points, -1, 1)

#compute the distance of each point from (0,0)
distance <-sqrt(x^2+y^2)

#boolean vector to indicate if each point is within the circle
within_circle <- ifelse(distance<1, TRUE, FALSE)
```

table(): Count the number of unique elements (TRUE and FALSE)

```
#compute proportion of points within circle/outside circle with table()
v <- table(within_circle)
```

```
#compute and print PI
pi <- v["TRUE"]/(v["TRUE"]+ v["FALSE"]) *4
print(pi)
```



output

TRUE
3.136

Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- **Monte Carlo Simulation with R**
 - Estimating π
 - **Simulating Product Demand**
 - Simulating Coin Flipping
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM

Simulating Product Demand

- Suppose the demand for a product is governed by the following discrete random variable:

Demand (X)	Probability
10,000	0.1
20,000	0.35
40,000	0.3
60,000	0.25

- How to simulate daily demand for 100 days and sum them?

Simulating Product Demand


- How to simulate daily demand for 100 days and sum them?

```
x <- c( 10000, 20000, 40000, 60000)
probability <- c(0.1, 0.35, 0.3, 0.25)

demand <- sample(x, 100, replace=TRUE, prob=probability)
barplot(table(demand))
```

Sample with replacement

Sum of the demand in 100 days



```
> sum(demand)
[1] 3470000
```

Expected value: $E[X] = 3500000$

Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- **Monte Carlo Simulation with R**
 - Estimating π
 - Simulating Product Demand
 - **Simulating Coin Flipping**
 - Stock Price Prediction
 - Hits on a Website
 - Queueing System for an ATM

Simulating Coin Flipping

- Given $n = 10$ flipping of a fair coin, what is the chance that we will get $k = 4$ heads-up or less?

```
n <- 10 # no. of coin flips
k <- 4  # no. of heads
runs <- 10000 # number of trials
```

```
# one trial simulates the flipping of a coin for 10 times
trial <- function(){
  sum(sample(c(0,1), n, replace=TRUE)) # of heads
}
```

One trial
experiment

```
#conduct trials 10000 times
result <- replicate(runs, trial())
t<-table(result)
print(t)
```

Output

## result	0	1	2	3	4	5	6	7	8	9	10
##	15	110	442	1167	2061	2465	2076	1118	460	73	13


Simulating Coin Flipping

- Given $n = 10$ flipping of a fair coin, what is the chance that we will get $k = 4$ heads-up or less?

```
#conduct trials 10000 times  
result <-replicate(runs, trial())  
t<-table(result)  
print(t)
```


```
barplot(table(result))
```

Draw the barplot to
visualize the sample results



Normalize

```
barplot(table(result)/runs)
```



Simulating Coin Flipping

- **Binomial Distribution**: The probability of getting k heads in n trials follows the binomial distribution.

- $\binom{n}{k} p^k p^{n-k}$

- *dbinom(x, size, prob)* : Probability density function

-

```
#prob(3 head out of 10 coin flips)
dbinom(3, size=10, prob=0.5)
```



```
## [1] 0.1171875
```

- *pbinom(q, size, prob, lower.tail)* : Calculate the cumulative probability

-

```
#prob(<=3 head out of 10 coin flips)
pbinom(3, size=10, prob=0.5, lower.tail= TRUE)
```

- **True:** $P(X \leq x)$
- **False:** $P(X > x)$

```
## [1] 0.171875
```

-

```
#prob(>3 head out of 10 coin flips)
pbinom(3, size=10, prob=0.5, lower.tail= FALSE)
```

```
## [1] 0.828125
```


Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- **Monte Carlo Simulation with R**
 - Estimating π
 - Simulating Product Demand
 - Simulating Coin Flipping
 - **Stock Price Prediction**
 - Hits on a Website
 - Queueing System for an ATM

Stock Price Prediction

- The price of a stock today is \$30
- In prior data , the price of the stock increases a factor of mean 1.001, with standard deviation of 0.005, i.e., price changes satisfy $N(1.001, 0.005^2)$
- What is the price of the stock after 365 days?


- ```
#generate the price change for 365 days
days <-365
changes <- rnorm(days, mean=1.001, sd=0.005)

price<-cumprod(c(30, changes)) #30*1st random no.*2nd random no...
Plot the line chart R Base Graphics
plot(price, type="l")
```

Returns a vector whose elements are the cumulative products of the elements of the argument.

# Stock Price Prediction

- The price of a stock today is \$30
- In prior data , the price of the stock increases a factor of mean 1.001, with standard deviation of 0.005, i.e., price changes satisfy  $N(1.001, 0.005^2)$
- What is the price of the stock after 365 days?



# Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte-Carlo Simulation
- **Monte-Carlo Simulation with R**
  - Estimating  $\pi$
  - Simulating Product Demand
  - Simulating Coin Flipping
  - Stock Price Prediction
  - **Hits on a Website**
  - Queueing System for an ATM

# Hits on a Website

- We are interested to estimate the number of hits on a certain website during a fixed time interval.
- The average rates per minute is 5
- **Poisson distribution** is often used to model rare events that are extremely unlikely to occur within a very short period of time or simultaneously (e.g. within 0.0001s).
  - **Examples:** arrivals of jobs to printers, arrivals of buses, occurrences of earthquake/traffic accidents/errors in software, arrival of customers to a queue.



# Hits on a Website

- **Poisson distribution** describes the probability of a given number of events occurring in a fixed interval of time and/or space
  - These events occur independently with a known average rate of the time since the last event
  - Poisson distribution has a single parameter  $\lambda > 0$ , being the average number of occurrence of the considered events.
  - $P(k \text{ events in an interval}) = e^{-\lambda} \frac{\lambda^k}{k!}$

Discrete or  
Continuous?



# Hits on a Website

- We are interested to estimate the number of hits on a certain website during a fixed time interval.
- The average rates per minute is 5
- **Task:** Find the probability that there will be exactly 17 hits in the next 3 minutes.

$$k = 17$$

Since the events occur independently,  $\lambda = 3 \times 5 = 15$ .





# Hits on a Website

- We are interested to estimate the number of hits on a certain website during a fixed time interval.
- The average rates per minute is 5
- **Task:** Find the probability that there will be exactly 17 hits in the next 3 minutes.

$$\lambda = 3 \times 5 = 15$$

- ***dpois(x, lambda)*** : Probability density function

```
> dpois(17, 15) #probability of 17 events with rate=15
[1] 0.08473555
```

- ***ppois(x, lambda, lower.tail)*** : Cumulative probability function


```
> ppois(17, 15, lower.tail=TRUE) #cumulative probability
[1] 0.7488588
```

# Hits on a Website

- We are interested to estimate the number of hits on a certain website during a fixed time interval.
- The average rates per minute is 5
- **Task:** Create a barplot to show the probability of having  *$n$  visitors in 1 minute*.

```
rate <- 5
x <- 1:20
p <- dpois(x, rate)
barplot(p, names.arg=x)
```

Output



How will the shape change with varying  $\lambda$ ?

# Hits on a Website

- What if we want to estimate the waiting time for the next hit?
- ***Exponential Distribution***: used to model the time that elapses before an event occurs, e.g., the time between two events.
- **Examples:**
  - Waiting time for the failure of a light bulb
  - Waiting time between customers coming into a store
  - Waiting time between phone calls to a hotline.
  - The length of a phone call

Discrete or  
Continuous?



# Hits on a Website

- What if we want to estimate the waiting time for the next hit?
- **Exponential Distribution**: used to model the time that elapses before an event occurs, e.g., the time between two events.
- Exponential Distribution VS. Poisson Distribution
  - The inter-arrival times of events in a Poisson process with rate  $\lambda$  is exponential and mean  $1/\lambda$ .
  - What is the average time between two hits (*5 hits/min*)?



# Hits on a Website

- What if we want to estimate the waiting time for the next hit?

Generates  $n$  random number from exponential distribution with certain rate  $\lambda$

```
#Waiting time for the next hit to the website
waitingTimeWebsite<-rexp(50, 5)
waitingTimeWebsite
```



```
[1] 0.09693668 0.14672284 0.04367192 0.03609941 0.06439198 0.39497901
[7] 0.05534583 0.45604551 0.32502948 0.15767379 0.07978570 0.68853209
[13] 0.60078705 0.27334693 0.06470556 0.05386009 0.06896579 0.03624342
[19] 0.26621996 0.62889533 0.06780421 0.15871193 0.15979114 0.21233654
[25] 0.47490147 0.01712544 0.10421384 1.02176814 0.06462670 0.06672129
[31] 0.19380114 0.03002739 0.26337777 0.27922920 0.04365630 0.02090912
[37] 0.03615915 0.59206056 0.05625460 0.00253461 0.17629357 0.81852976
[43] 0.00953266 0.42251408 0.09151426 0.65567584 0.35261606 0.37254870
[49] 0.08170725 0.62902729
```

How to  
visualize it for  
result analysis?

1. Barplot
2. Histogram
3. Scatterplot


# Hits on a Website

- What if we want to estimate the waiting time for the next hit?

Assign the elements into 20 bins

```
hist(waitingTimeWebsite, breaks=20)
```

**Histogram of waitingTimeWebsite**




# Roadmap

- Introduction: *History* and *Motivation*
- Key Factors in Monte Carlo Simulation
- **Monte Carlo Simulation with R**
  - Estimating  $\pi$
  - Simulating Product Demand
  - Simulating Coin Flipping
  - Stock Price Prediction
  - Hits on a Website
  - **Queueing System for an ATM**

# Queueing System for ATM

- We will model a simple queueing system that represents customers arriving at and using an ATM.
- Customers arrive at an ATM and perform a number of ATM transactions.
- One customer can use ATM at a time. After completing their transactions, the customer takes their card and leaves.
- When the ATM is in use, the next customers have to wait until the previous customer has left.





# Queueing System for an ATM

- The first customer arrives at the ATM at 1 min and takes 1.3 minutes to complete his transactions.
- The second customer arrives at time 2 min and takes 0.5 minutes to complete his transaction.
- Quick Questions:
  - At what time will the first customer leave?
  - How long will the second customer wait before he can use the ATM?
  - At what time will the second customer leave?

# Queueing System for an ATM

- A customer's waiting time depends on his arrival time and the time at which the previous customer will leave the ATM.
- There is no need for the first customer to wait.
- When determining if customer  $i$  should wait, we need to check when customer  $i - 1$  have completed his transactions when customer  $i$  arrives.



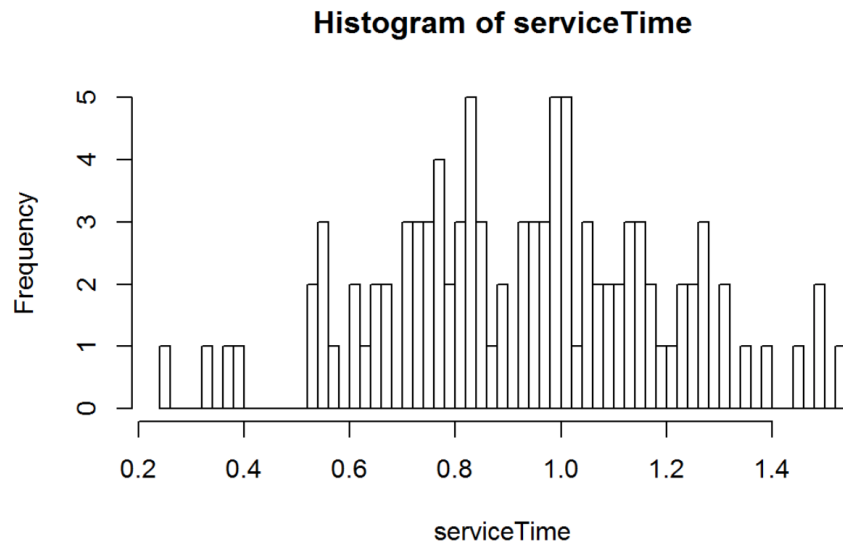

# Queueing System for an ATM

- Suppose that we want to **simulate** *100 customers* coming to the ATM.
- The **service time** follows a normal distribution with a mean of 0.9 minutes and standard deviation of 0.25 minutes  $N(0.9, 0.25^2)$ .
- The **inter-arrival time of customers** follows an *exponential distribution* with a rate of 1 customer per minute ( $\lambda = 1$ )
- How to do the simulation?

# Queueing System for an ATM

- Service times can be sampled from the normal distribution using the *`rnorm(100, mean=0.9, sd=0.25)`* function

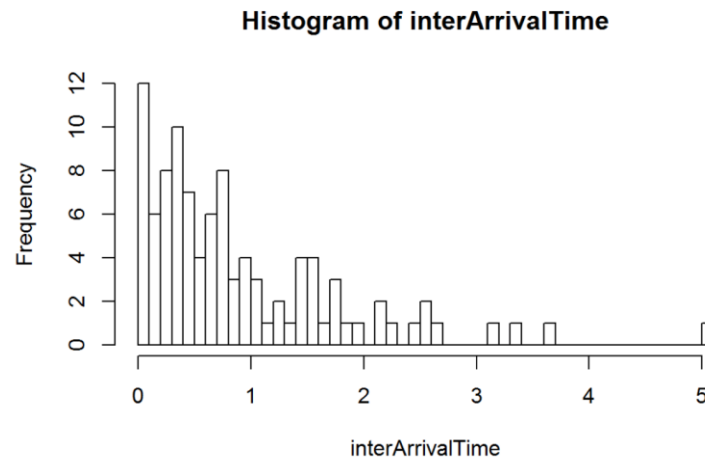

```
serviceTime<-rnorm(100, 0.9, 0.25)
hist(serviceTime, breaks=50)
```



# Queueing System for an ATM

- The interarrival time of the customers can be sampled using `rexp(100, rate=1)`.
- This may simulate the time for customer arrived after the previous customer.


```
interArrivalTime<-rexp(100, 1)
hist(interArrivalTime, breaks=50)
```



# Queueing System for an ATM

- Use the function *cumsum()* to compute the actual arrival time for the 100 customers.

```
arrivalTime<-cumsum(rexp(100, 1))
arrivalTime
```



In increasing order

```
> arrivalTime
[1] 0.9760055 1.0091946 1.0889435 1.3630373 1.4842585 3.7292324
[7] 5.4970826 6.0633743 6.3316084 6.5905567 10.4530757 11.9483586
[13] 12.1467582 12.2360948 14.1136987 15.0155603 15.5390824 17.2292949
[19] 18.2237556 19.3294594 20.2104632 20.3775642 21.1634930 24.2176364
[25] 24.3553744 25.1767660 25.3426509 25.9323925 27.0038995 29.0078009
[31] 29.2887315 29.4327757 32.2866256 32.3242688 32.3741418 34.1917064
[37] 34.2460832 34.2564507 39.3157243 40.5592686 42.4860609 43.8888613
[43] 45.7021042 46.8687582 48.4617725 48.6487968 49.5481331 49.7236396
[49] 49.9856995 50.9172966 51.0778695 51.5505652 52.9261202 53.7540612
[55] 54.3703905 55.7960763 56.7781005 59.3511263 60.0448332 61.6198472
[61] 62.4108903 64.2093477 65.7146795 67.1512490 68.1908882 68.8950934
[67] 69.5691299 70.1948541 71.4499252 72.4913475 73.5774646 75.1575037
[73] 76.0861064 77.7882944 78.3753304 78.3861295 81.5810221 82.2756851
[79] 83.7405187 86.2334013 86.4867768 86.5845150 87.4119343 91.2065208
[85] 92.5837602 92.8227321 93.6039122 94.3334029 102.6586363 104.0988170
[91] 105.2980914 105.3457986 107.3845934 107.4975768 108.4143292 108.7886023
[97] 108.8879160 109.3589560 110.4997033 110.6833009
```

# Queueing System for an ATM

- Let  $a_i$  and  $s_i$  be the arrival time and service time for customer  $i$ .
- Also, let  $l_{i-1}$  be the leaving time of customer  $i - 1$ .
- Then, the waiting time for customer  $i$  is computed as follows:
  - If  $a_i > l_{i-1}$ , customer  $i$  will arrive after customer  $i - 1$  has left. Then, the waiting time for customer  $i$  is  $w_i = 0$ .
  - Else,  $w_i = l_{i-1} - a_i$ .
- So,  $w_i = \max(0, l_{i-1} - a_i)$
- The customer  $i$  will leave at  $l_i = a_i + w_i + s_i$



# A slide to takeaway

- **Concept:** What is Monte Carlo Simulation?
- **Methods:** How to design Monte Carlo simulation methods to estimate a model (a probability distribution)?
- **Implementation:** How to use R programming to implement Monte Carlo simulation?